

Assignment 1)Write a python program to Prepare Scatter Plot (Use Forge Dataset / Iris Dataset)

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
```

Step 1: Load the Iris dataset

```
column_names = ['sepal_length', 'sepal_width', 'petal_length',
                'petal_width']
iris=load_iris()
df=pd.DataFrame(data=iris.data,columns=column_names)
df['species']=iris.target
```

Step 2: Display the first 5 rows

```
print("First 5 rows of the Iris dataset:")
print(df.head())
```

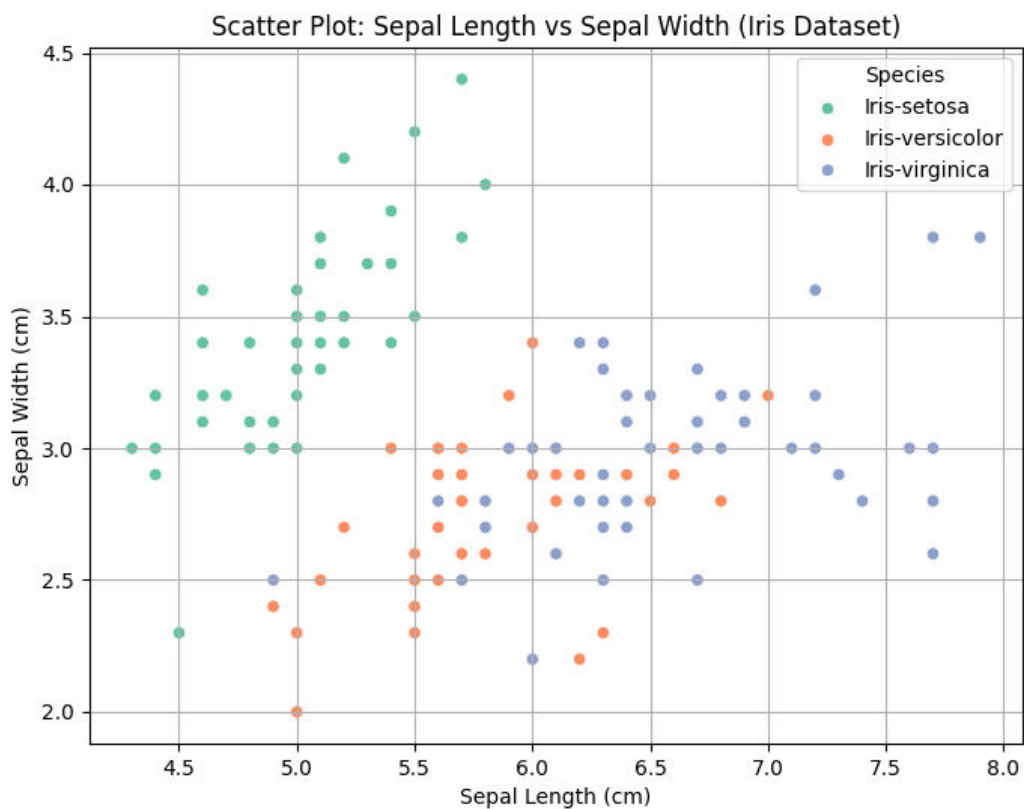
Step 3: Create a scatter plot: Sepal Length vs Sepal Width

```
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x='sepal_length', y='sepal_width',
                hue='species', palette='Set2')
plt.title('Scatter Plot: Sepal Length vs Sepal Width (Iris Dataset)')
plt.xlabel('Sepal Length (cm)')
plt.ylabel('Sepal Width (cm)')
plt.legend(title='Species')
plt.grid(True)
plt.show()
```

output:

First 5 rows of the Iris dataset:

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa



**Assignment 2)Write a python program to find all null values in a given data set and remove them.
Create your own dataset.**

```
import pandas as pd
import numpy as np
```

Creating a sample DataFrame with missing values

```
data = {
    'School ID': [101, 102, 103, np.nan, 105, 106, 107, 108],
    'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eva', 'Frank', 'Grace',
'Henry'],
    'Address': ['123 Main St', '456 Oak Ave', '789 Pine Ln', '101 Elm St',
np.nan, '222 Maple Rd', '444 Cedar Blvd', '555 Birch Dr'],
    'City': ['Los Angeles', 'New York', 'Houston', 'Los Angeles', 'Miami',
np.nan, 'Houston', 'New York'],
    'Subject': ['Math', 'English', 'Science', 'Math', 'History', 'Math',
'Science', 'English'],
    'Marks': [85, 92, 78, 89, np.nan, 95, 80, 88],
    'Rank': [2, 1, 4, 3, 8, 1, 5, 3],
    'Grade': ['B', 'A', 'C', 'B', 'D', 'A', 'C', 'B']
}
```

```
df = pd.DataFrame(data)
print("Sample DataFrame:")
print(df)
```

```
df_cleaned = df.dropna() #this method is used to remove null values  
from given dataset
```

Displaying the DataFrame after removing missing values

```
print("\nDataFrame after removing rows with missing values:")  
print(df_cleaned)
```

output:

Sample DataFrame:

	School ID	Name	Address	City	Subject	Marks	Rank	Grade
0	101.0	Alice	123 Main St	Los Angeles	Math	85.0	2	B
1	102.0	Bob	456 Oak Ave	New York	English	92.0	1	A
2	103.0	Charlie	789 Pine Ln	Houston	Science	78.0	4	C
3	NaN	David	101 Elm St	Los Angeles	Math	89.0	3	B
4	105.0	Eva	NaN	Miami	History	NaN	8	D
5	106.0	Frank	222 Maple Rd	NaN	Math	95.0	1	A
6	107.0	Grace	444 Cedar Blvd	Houston	Science	80.0	5	C
7	108.0	Henry	555 Birch Dr	New York	English	88.0	3	B

DataFrame after removing rows with missing values:

	School ID	Name	Address	City	Subject	Marks	Rank	Grade
0	101.0	Alice	123 Main St	Los Angeles	Math	85.0	2	B
1	102.0	Bob	456 Oak Ave	New York	English	92.0	1	A
2	103.0	Charlie	789 Pine Ln	Houston	Science	78.0	4	C
6	107.0	Grace	444 Cedar Blvd	Houston	Science	80.0	5	C
7	108.0	Henry	555 Birch Dr	New York	English	88.0	3	B

Assignment 3) Write a python program the Categorical values in numeric format for a given dataset.(iris dataset)

```
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.preprocessing import LabelEncoder
```

Step 1: Load the Iris dataset

```
iris = load_iris()
df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
```

Step 2: Add the species column using human-readable names

```
df['species'] = [iris.target_names[i] for i in iris.target]
```

Step 3: Label Encode the 'species' column

```
le = LabelEncoder()
df['species_encoded'] = le.fit_transform(df['species'])
```

Step 4: Display the original and encoded data

```
print("Original DataFrame with Categorical 'species':")
print(df[['species']].head())
```

```
print("\nDataFrame after Label Encoding:")
print(df[['species', 'species_encoded']].head())
```

Step 5: Show label mapping

```
print("\nLabel Encoding Mapping:")
for class_label, encoded_val in zip(le.classes_,
le.transform(le.classes_)):
```

```
print(f"{class_label} → {encoded_val}")
```

output:

Original DataFrame with Categorical 'species':

```
species
0 setosa
1 setosa
2 setosa
3 setosa
4 setosa
```

DataFrame after Label Encoding:

	species	species_encoded
0	setosa	0
1	setosa	0
2	setosa	0
3	setosa	0
4	setosa	0

Label Encoding Mapping:

setosa → 0

versicolor → 1

virginica → 2

Assignment 4)Write a python program to implement simple Linear Regression for predicting house price .

```
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import r2_score
```

1. Manually creating the dataset

```
data = {  
    'Area': [1000, 1500, 1800, 2400, 3000, 3500],  
    'Price': [150000, 200000, 230000, 300000, 360000, 400000]  
}  
df = pd.DataFrame(data)
```

2. Define features and target

```
X = df[['Area']]  
y = df['Price']
```

3. Split the data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)
```

4. Train the linear regression model

```
model = LinearRegression()  
model.fit(X_train, y_train)
```

5. Predict price for 2000 sq. ft.

```
predicted_price = model.predict(pd.DataFrame([[2000]],  
columns=['Area']))  
print("Predicted Price for 2000 sq. ft.:", predicted_price[0])
```

6. Predict on the test set and compute R^2

```
y_pred = model.predict(X_test)  
r2 = r2_score(y_test, y_pred)  
print("R-squared ( $R^2$ ) Score:", r2)
```

7. Plot the training data, regression line, and test data

```
plt.scatter(X_train, y_train, color='blue', label='Training Data')  
plt.scatter(X_test, y_test, color='green', label='Test Data')  
plt.plot(X, model.predict(X), color='red', label='Regression Line')
```

```
plt.xlabel('Area (sq. ft.)')
plt.ylabel('Price')
plt.title('House Price vs Area')
plt.legend()
plt.grid(True)
plt.show()
```

OUTPUT:

Predicted Price for 2000 sq. ft.: 254792.6267281106

R-squared (R^2) Score: 0.9666885165443214



**Assignment 5) Write a python program to implement multiple Linear Regression for a given dataset.
(Use 50Startups.csv dataset)**

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.metrics import r2_score
```

Step 1: Load the dataset

```
dataset = pd.read_csv("D:/Teaching Material/Machine
Learning/Assignment 5/50_Startups.csv")
print(dataset.head())
```

Step 2: Separate features (X) and target (y)

```
X = dataset.iloc[:, :-1].values # All columns except the last
y = dataset.iloc[:, -1].values # Only the last column (Profit)
```

Step 3: Encode categorical data (State column)

```
ct = ColumnTransformer(
    transformers=[('encoder', OneHotEncoder(), [3])], # Column 3 is
    'State'
    remainder='passthrough'
)
```

```
X = np.array(ct.fit_transform(X)) # Convert to NumPy array after encoding
```

Step 4: Split the dataset into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)
```

Step 5: Train the Multiple Linear Regression model

```
regressor = LinearRegression()  
regressor.fit(X_train, y_train)
```

Step 6: Predict the results on the test set

```
y_pred = regressor.predict(X_test)
```

Step 7: Compare predicted and actual values

```
comparison_df = pd.DataFrame({  
    'Actual Profit': y_test,  
    'Predicted Profit': y_pred  
})  
print(comparison_df)
```

#Step 8: Evaluate model performance using R^2 score

```
print("\n $R^2$  Score:", r2_score(y_test, y_pred))
```

Optional: Plot predicted vs actual


```
plt.scatter(y_test, y_pred, color='blue', alpha=0.6)  
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],  
color='red')  
plt.xlabel('Actual Profit')  
plt.ylabel('Predicted Profit')  
plt.title('Actual vs Predicted Profit')  
plt.grid(True)  
plt.tight_layout()  
plt.show()
```

output:

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

	Actual Profit	Predicted Profit
0	134307.35	126362.879083
1	81005.76	84608.453836
2	99937.59	99677.494252
3	64926.08	46357.460686
4	125370.37	128750.482885
5	35673.41	50912.417419
6	105733.54	109741.350327
7	107404.34	100643.242816
8	97427.84	97599.275746
9	122776.86	113097.425244

R² Score: 0.8987266414319911



Assignment 6) Write a python program to implement Polynomial Regression for given dataset. (use salary_positions.csv dataset)

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score
```

Load dataset

```
data = pd.read_csv('D:\\Teaching Material\\Machine
Learning\\Assignment 6 Polynomial
Regression\\salary_positions.csv')
X = data[['Level']].values
y = data['Salary'].values
```

Split dataset

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Polynomial features (degree 4)

```
poly = PolynomialFeatures(degree=4)
X_train_poly = poly.fit_transform(X_train)
X_test_poly = poly.transform(X_test)
```

Train model

```
model = LinearRegression()
model.fit(X_train_poly, y_train)
```

Predict on test and all data

```
y_test_pred = model.predict(X_test_poly)
y_all_pred = model.predict(poly.transform(X))
```

Evaluation

```
print("R2 Score:", r2_score(y_test, y_test_pred))
```

Predict for Level 11 and 12

```
salary_11 = model.predict(poly.transform([[11]]))[0]
salary_12 = model.predict(poly.transform([[12]]))[0]
print("Predicted Salary for Level 11:", round(salary_11, 2))
print("Predicted Salary for Level 12:", round(salary_12, 2))
```

Plot without smoothing

```
plt.scatter(X, y, color='blue', label='Actual Salaries')
plt.plot(X, y_all_pred, color='red', marker='o', label='Predicted Curve (No Smoothing)')
plt.title('Polynomial Regression (No Smoothing)')
plt.xlabel('Position Level')
plt.ylabel('Salary')
plt.legend()
plt.grid(True)
```

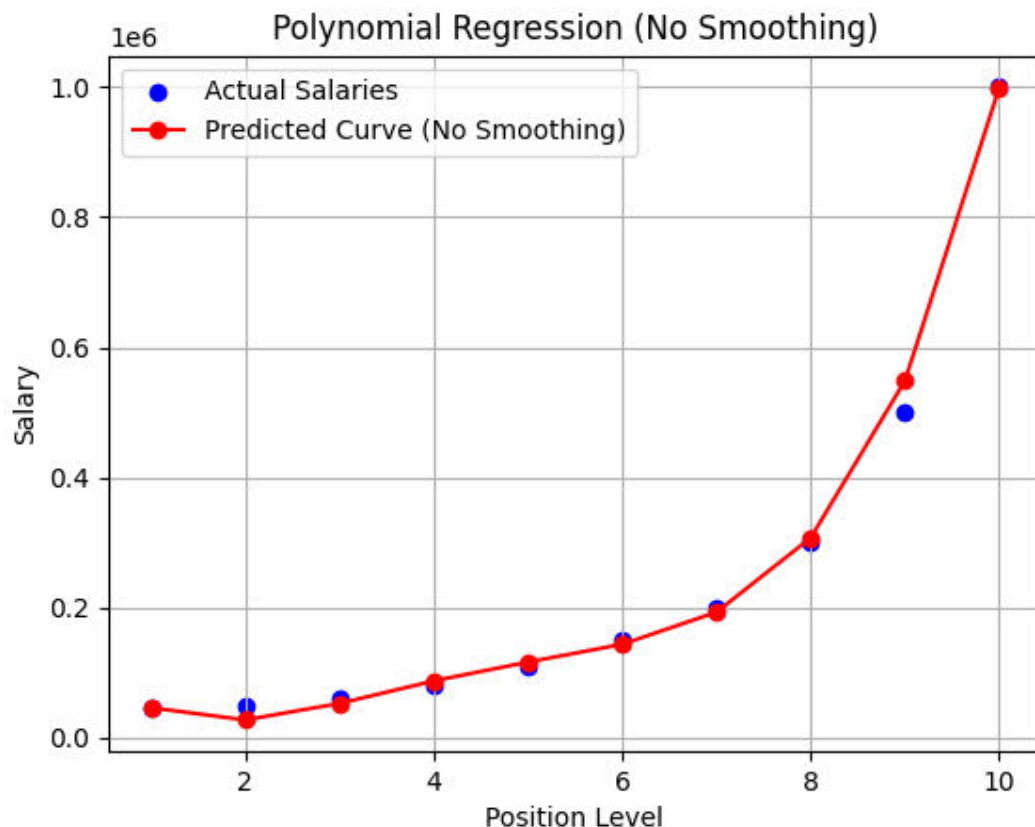
```
plt.show()
```

Output:

R^2 Score: 0.9714666803843249

Predicted Salary for Level 11: 1759103.68

Predicted Salary for Level 12: 2949328.94



Assignment 7(a): Write a python program to implement Logistic Regression for Iris dataset

Logistic Regression on Iris Dataset

```
from sklearn.datasets import load_iris
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
```

1. Load the dataset

```
iris = load_iris()
X = iris.data # Features
y = iris.target # Target labels
```

2. Split into training and testing sets (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

3. Create and train Logistic Regression model

```
model = LogisticRegression(max_iter=200)
model.fit(X_train, y_train)
```

4. Make predictions

```
y_pred = model.predict(X_test)
```

5. Evaluate accuracy

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Print classification report

```
print("\nClassification Report:\n")
print(classification_report(y_test, y_pred,
    target_names=iris.target_names))
```

6. Test prediction on a sample

```
sample = [[5.1, 3.5, 1.4, 0.2]] # Example flower features
pred_class = iris.target_names[model.predict(sample)[0]]
print("\nSample Prediction:", pred_class)
```

Output:

Accuracy: 1.0

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Sample Prediction: setosa

Assignment 7(b): Write a python program to Implement Naïve Bayes.

Import necessary libraries

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report
```

Load dataset

```
columns = ["Pregnancies", "Glucose", "BloodPressure",
           "SkinThickness",
           "Insulin", "BMI", "DiabetesPedigreeFunction", "Age",
           "Outcome"]
```



```
data = pd.read_csv("D:\Teaching Material\Machine Learning\Assignment7(b)\diabetes_dataset.csv",header=None,names=columns)
```

Show first 5 rows

```
print("First 5 rows of dataset:")  
print(data.head())
```

```
X = data.drop("Outcome", axis=1) # Drop Outcome column → features  
y = data["Outcome"]
```

Split into train and test sets (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)
```

Initialize Gaussian Naive Bayes model

```
model = GaussianNB()
```

Train the model

```
model.fit(X_train, y_train)
```

Make predictions

```
y_pred = model.predict(X_test)
```

Evaluate the model

```
print("Accuracy:", accuracy_score(y_test, y_pred))  
print("\nClassification Report:\n", classification_report(y_test,  
y_pred))
```

Sample prediction:

```
# Format: [Pregnancies, Glucose, BloodPressure, SkinThickness,  
Insulin, BMI, DiabetesPedigreeFunction, Age]
```

```
sample = [[6,148,72,35,0,33.6,0.627,50]]
```

```
prediction = model.predict(sample)
```

```
if prediction[0] == 1:
```

```
    print("\nThe sample input is predicted to HAVE diabetes.")
```

```
else:
```

```
    print("\nThe sample input is predicted to NOT have diabetes.")
```

OUTPUT:

First 5 rows of dataset:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Accuracy: 0.7662337662337663

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.80	0.81	99
1	0.66	0.71	0.68	55
accuracy			0.77	154
macro avg	0.75	0.75	0.75	154
weighted avg	0.77	0.77	0.77	154

The sample input is predicted to HAVE diabetes.

Assignment 8: Write a python program to Implement Decision Tree whether or not to play tennis

Step 1: Import libraries

```
import pandas as pd
from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
```

Step 2: Load dataset from local CSV file

```
df = pd.read_csv("D:\Teaching Material\Machine Learning\Assignment 8\Play_tennis.csv")
```

Step 3: Encode categorical variables

```
le = LabelEncoder()
for col in df.columns:
    df[col] = le.fit_transform(df[col])
```

Step 4: Split features (X) and target (y)

```
X = df[['Outlook','Temperature','Humidity','Wind']]
y = df['Play Tennis']
```

Step 5: Train Decision Tree

```
model = DecisionTreeClassifier(criterion='entropy')
```

```
model.fit(X, y)
```

Step 6: Visualize Decision Tree

```
plt.figure(figsize=(10,6))
```

```
plot_tree(model,
```

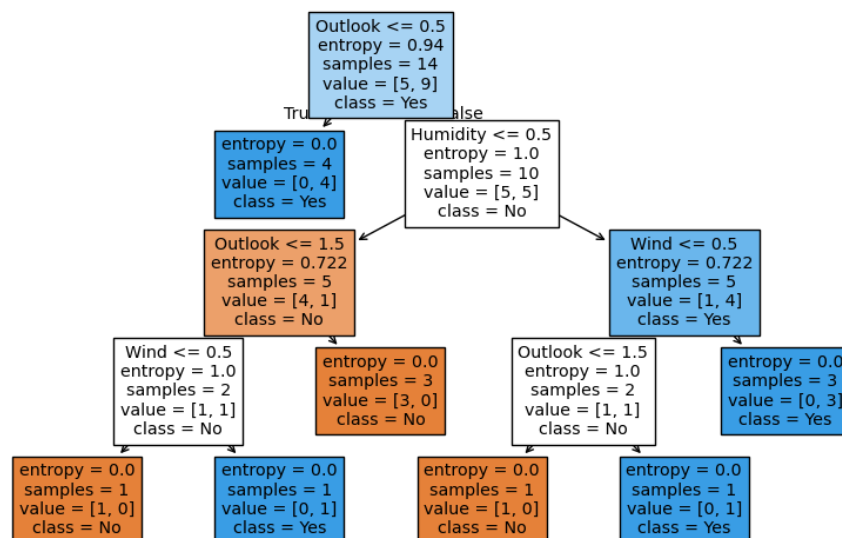
```
    feature_names=['Outlook','Temperature','Humidity','Wind'],
```

```
    class_names=['No','Yes'],
```

```
    filled=True)
```

```
plt.show()
```

OUTPUT:



Assignment(9): Write a python program to Implement Random Forest for iris Dataset

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix
```

1. Load dataset

```
iris = load_iris()
X = iris.data
y = iris.target
```

2. Split dataset

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
random_state=42)
```

3. Train Random Forest model

```
model = RandomForestClassifier(n_estimators=100,
random_state=42)
model.fit(X_train, y_train)
```

4. Predictions

```
y_pred = model.predict(X_test)
```

5. Accuracy

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

6. Confusion Matrix

```
cm = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:\n", cm)
```

7. Prediction for a new flower

```
sample = [[5.0, 3.6, 1.4, 0.2]]
print("\nPrediction for sample:",
iris.target_names[model.predict(sample)][0])
```

output:

Accuracy: 1.0

Confusion Matrix:

```
[[19  0  0]
```

```
 [ 0 13  0]
```

```
 [ 0  0 13]]
```

Prediction for sample: setosa

**Assignment (10): Write a python program to implement linear SVM.
(use wine dataset)**

```
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix
```

1. Load dataset

```
wine = datasets.load_wine()
X = wine.data
y = wine.target
```

2. Split dataset

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)
```

3. Train Linear SVM

```
model = SVC(kernel='linear', C=1.0, random_state=42)
model.fit(X_train, y_train)
```

4. Predictions

```
y_pred = model.predict(X_test)
```

5. Evaluation

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

```
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

6. Example prediction

```
sample = [X_test[0]]  
print("\nPrediction for sample:",  
wine.target_names[model.predict(sample)][0])
```

OUTPUT:

Accuracy: 0.9814814814814815

Confusion Matrix:

```
[[19  0  0]  
 [ 0 20  1]  
 [ 0  0 14]]
```

Prediction for sample: class_0

Assignment(11): Write a python program to find Decision boundary by using a neural network with 10 hidden units on two moons dataset

```
import numpy as np  
import matplotlib.pyplot as plt  
from sklearn.datasets import make_moons  
from sklearn.model_selection import train_test_split  
from sklearn.neural_network import MLPClassifier
```


1. Generate two moons dataset

```
X, y = make_moons(n_samples=1000, noise=0.2, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)
```

2. Define and train Neural Network with 10 hidden units

```
model = MLPClassifier(hidden_layer_sizes=(10,), activation='relu',
solver='adam', max_iter=2000, random_state=42)
model.fit(X_train, y_train)
```

3. Plot decision boundary

```
def plot_decision_boundary(model, X, y):
    x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
    y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5
    h = 0.01
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                          np.arange(y_min, y_max, h))


    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    plt.contourf(xx, yy, Z, alpha=0.4)
    plt.scatter(X[:, 0], X[:, 1], c=y, s=20, edgecolor='k')
```

```
plt.show()
```

```
plot_decision_boundary(model, X, y)
```

OUTPUT:



Assignment(12): Write a python program to generate frequent itemset and association rule by applying apriori algorithm on Market basket dataset

```
import pandas as pd
```

```
from mlxtend.frequent_patterns import apriori, association_rules
```

```
from mlxtend.preprocessing import TransactionEncoder
```

Sample Market Basket Dataset (replace with your actual data)

```
dataset = [['Milk', 'Onion', 'Nutmeg', 'Kidney Beans', 'Eggs', 'Yogurt'],  
           ['Dill', 'Onion', 'Nutmeg', 'Kidney Beans', 'Eggs', 'Yogurt'],  
           ['Milk', 'Apple', 'Kidney Beans', 'Eggs'],  
           ['Milk', 'Unicorn', 'Corn', 'Kidney Beans', 'Yogurt'],  
           ['Corn', 'Onion', 'Onion', 'Kidney Beans', 'Ice cream', 'Eggs']]
```

Convert the dataset into a one-hot encoded DataFrame

```
te = TransactionEncoder()  
te_ary = te.fit(dataset).transform(dataset)  
df = pd.DataFrame(te_ary, columns=te.columns_)
```

Generate frequent itemsets using Apriori

```
frequent_itemsets = apriori(df, min_support=0.6,  
                             use_colnames=True)
```

Generate association rules

```
rules = association_rules(frequent_itemsets, metric="confidence",  
                          min_threshold=0.7, num_itemsets=5)
```

Display frequent itemsets

```
print("Frequent Itemsets:")
```

```
print(frequent_itemsets)
```

Display association rules

```
print("\nAssociation Rules:")
```

```
print(rules[['antecedents', 'consequents', 'support', 'confidence',  
'lift']])
```

OUTPUT:

Frequent Itemsets:

	support	itemsets
0	0.8	(Eggs)
1	1.0	(Kidney Beans)
2	0.6	(Milk)
3	0.6	(Onion)
4	0.6	(Yogurt)
5	0.8	(Kidney Beans, Eggs)
6	0.6	(Eggs, Onion)
7	0.6	(Milk, Kidney Beans)
8	0.6	(Kidney Beans, Onion)
9	0.6	(Kidney Beans, Yogurt)
10	0.6	(Kidney Beans, Eggs, Onion)

Association Rules:

	antecedents	consequents	support	confidence	lift
0	(Kidney Beans)	(Eggs)	0.8	0.80	1.00

1	(Eggs)	(Kidney Beans)	0.8	1.00	1.00
2	(Eggs)	(Onion)	0.6	0.75	1.25
3	(Onion)	(Eggs)	0.6	1.00	1.25
4	(Milk)	(Kidney Beans)	0.6	1.00	1.00
5	(Onion)	(Kidney Beans)	0.6	1.00	1.00
6	(Yogurt)	(Kidney Beans)	0.6	1.00	1.00
7	(Kidney Beans, Eggs)	(Onion)	0.6	0.75	1.25
8	(Kidney Beans, Onion)	(Eggs)	0.6	1.00	1.25
9	(Eggs, Onion)	(Kidney Beans)	0.6	1.00	1.00
10	(Eggs)	(Kidney Beans, Onion)	0.6	0.75	1.25
11	(Onion)	(Kidney Beans, Eggs)	0.6	1.00	1.25

Assignment(13): Write a python program to implement k-nearest Neighbours ML algorithm to build prediction model (Use Forge Dataset)

```
import matplotlib.pyplot as plt
```

```
import mglearn
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import accuracy_score, confusion_matrix
```

1. Load the Forge dataset

```
X, y = mglearn.datasets.make_forge()
```

2. Split dataset into training and testing

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
random_state=42)
```

3. Create KNN model (k=3)

```
knn = KNeighborsClassifier(n_neighbors=3)  
knn.fit(X_train, y_train)
```

4. Predictions

```
y_pred = knn.predict(X_test)
```

5. Accuracy

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

6. Confusion Matrix

```
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

7. Visualization of decision boundary

```
mglearn.plots.plot_2d_separator(knn, X, fill=True, eps=0.5,  
alpha=0.5)
```

```
mglearn.discrete_scatter(X[:, 0], X[:, 1], y)
```

```
plt.legend(["Class 0", "Class 1"], loc=4)
```

```
plt.xlabel("Feature 1")
```

```
plt.ylabel("Feature 2")  
plt.title("KNN Decision Boundary (k=3)")  
plt.show()
```


OUTPUT:

Accuracy: 0.8571428571428571

Confusion Matrix:

```
[[1 0]
```

```
[1 5]]
```



Assignment(14): Write a python program to implement Agglomerative clustering on a synthetic dataset.

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_blobs
from sklearn.cluster import AgglomerativeClustering
```

Step 1: Generate synthetic dataset

```
X, y_true = make_blobs(n_samples=200, centers=4, cluster_std=1.0,
random_state=42)
```

Step 2: Apply Agglomerative Clustering

```
agg_clust = AgglomerativeClustering(n_clusters=4, linkage='ward')
y_pred = agg_clust.fit_predict(X)
```

Step 3: Plot the clusters

```
plt.figure(figsize=(8,6))
plt.scatter(X[:, 0], X[:, 1], c=y_pred, cmap='rainbow', s=50)
plt.title("Agglomerative Clustering on Synthetic Dataset")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.show()
```

OUTPUT:

